





# Tennis Shots Classification With Deep Learning

MSCA 37011 Spring 2022

Vritti Gandhi, Shikhar Madrecha, Haider Waseem, Nihaal Zaveri

# Agenda

-  Abstract
-  Exploratory Data Analysis
-  Modeling
-  Results and Future Work

# Abstract

- 🎾 Goal: identify the type of tennis shots from a video
- 🎾 Data taken from Thetis dataset
- 🎾 Implemented three different models
  - Base CNN model
  - Binary classification CNN model
  - LRCNN


EDA

# Data Dictionary

The data has been taken from ***THree DimEnsional TennIs Shots (THETIS)*** dataset.

THETIS set is comprised of a set of 12 basic Tennis shots performed by 31 amateurs and 24 experienced players. Each shot has been performed several times resulting in 1980 (single period cropped) videos, converted to AVI format. Videos were captured using an Xbox Kinect.

The total duration of the videos is 2 hours and 15 minutes.



Two handed backhand



Forehand Flat



One handed backhand




Smash

The complete set of classes which are included in the dataset are:

1. *backhand,*
2. *backhand 2 hands,*
3. *backhand slice,*
4. *backhand volley,*
5. *flat service,*
6. *forehand flat,*
7. *forehand open stands,*
8. *forehand slice,*
9. *forehand volley,*
10. *kick service,*
11. *slice service,*
12. *smash*

# Distribution of Classes and Data Size



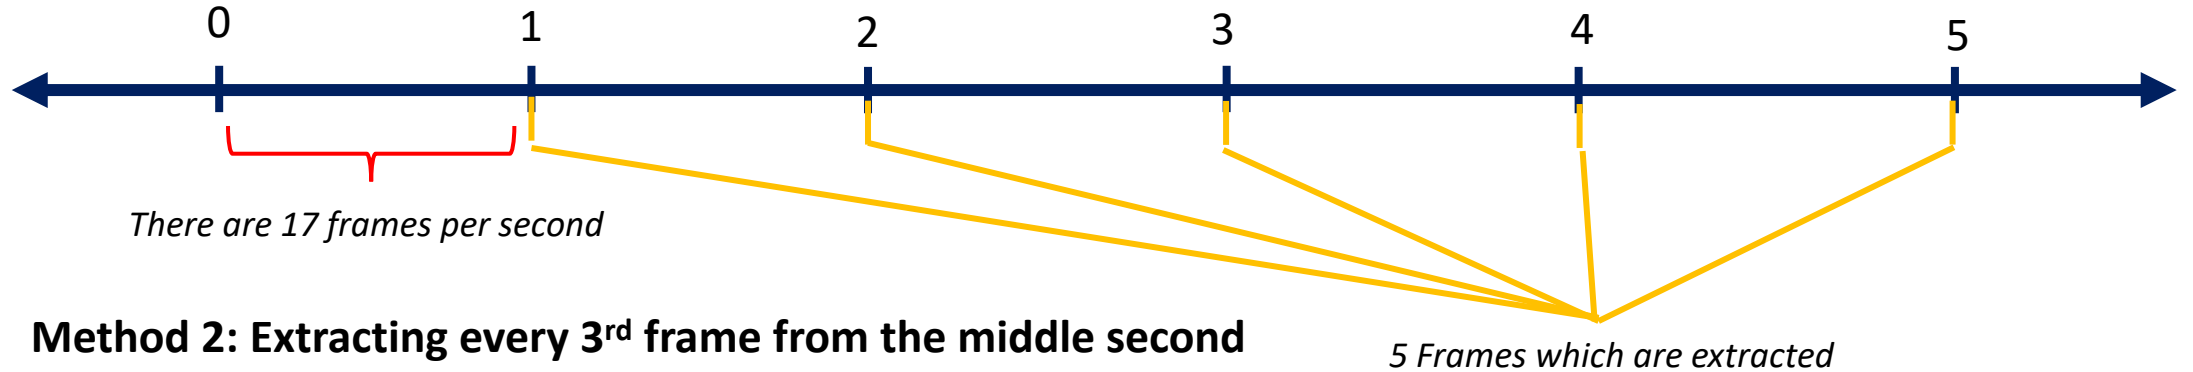
The THETIS dataset contains 5 second videos for each shot. To be able to input these videos in the model, we split every video in different frames and use these frames as an input into the model.

We have total **1980** videos split equally between 12 classes, i.e., **165 videos for each class**.


From these videos, we extract the frame at the end of every second which leads to generation of c. **9k frames** in total and around **700-800 frames for each class**.

# Extracting Frames


## Method 1: Extracting the last frame of every second



## Method 2: Extracting every 3<sup>rd</sup> frame from the middle second



## Method 3: Extracting 16 frames from the entire video



# Resizing Frames

The original size of the frames extracted from videos was 480\*640 pixels. This was too large to be used as an input on our local machines and was making the model run into OOM errors.

We decided to resize the images to 120\*160 pixels to improve model speed and performance.

We also normalize all images by dividing the pixels by 255.





# Infrastructure Used

As running CNNs can be computationally very expensive, we have tried variety of platforms / infrastructures to boost training and prediction speed.



# Modeling

# Base Model

🟡 For our baseline model, we use all 12 shot classes

🟡 We selected 200 frames of each shot type for training the model

🟡 50 frames of each shot were used for validation


Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 118, 158, 32)	896
conv2d_1 (Conv2D)	(None, 114, 154, 64)	51264
batch_normalization (Batch Normalization)	(None, 114, 154, 64)	256
conv2d_2 (Conv2D)	(None, 114, 154, 64)	102464
max_pooling2d (MaxPooling2D)	(None, 57, 77, 64)	0
batch_normalization_1 (Batch Normalization)	(None, 57, 77, 64)	256
flatten (Flatten)	(None, 280896)	0
dense (Dense)	(None, 12)	3370764
=====		
Total params: 3,525,900		
Trainable params: 3,525,644		
Non-trainable params: 256		

# Base Model



Test loss: [36.97636795043945, 0.10833333432674408]  
 Train loss: [16.63172721862793, 0.2212500125169754]



# Binary Classification

🎾 Instead of using all 12 shot classes, we simplified our approach to a binary classification

🎾 Groundstrokes: 2 handed backhand and flat forehand

🎾 Service: combine three types of serve – slice, flat, kick

🎾 1776 groundstroke frames


🎾 2553 service frames

Model: "sequential"

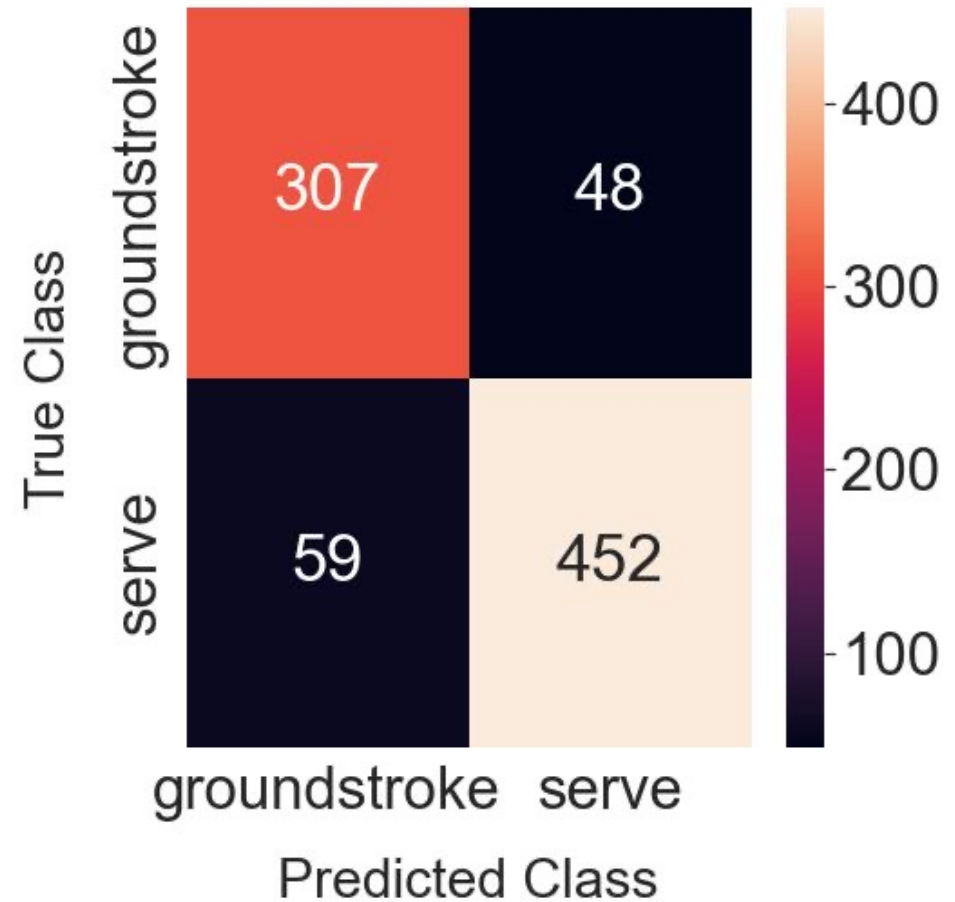
Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 118, 158, 32)	896
conv2d_1 (Conv2D)	(None, 116, 156, 64)	18496
batch_normalization (Batch Normalization)	(None, 116, 156, 64)	256
conv2d_2 (Conv2D)	(None, 116, 156, 64)	65600
max_pooling2d (MaxPooling2D)	(None, 58, 78, 64)	0
batch_normalization_1 (Batch Normalization)	(None, 58, 78, 64)	256
flatten (Flatten)	(None, 289536)	0
dense (Dense)	(None, 2)	579074

=====  
Total params: 664,578  
Trainable params: 664,322  
Non-trainable params: 256  
=====


# Binary Classification



Test loss: [6.450009822845459, 0.8764433860778809]  
Train loss: [3.3866405487060547, 0.9347386956214905]



# Binary Classification – 4D Input



One input

Instead of using each frame as a single data point, we also decided to use 4 frames from each video as a single data point and created a 4-dimensional input array instead of a 3-dimension input array.

```
[ ] x_train_resized.shape
    (2304, 120, 160, 3)

[ ] x_train_vids.shape
    (576, 4, 120, 160, 3)
```

Old input array (3d)

New input array (4d)

The idea behind this approach was to somehow provide a sequence of inputs in the traditional CNN architecture.

The **accuracy** from the binary classification problem is **around 60%** from this model.

Model Architecture

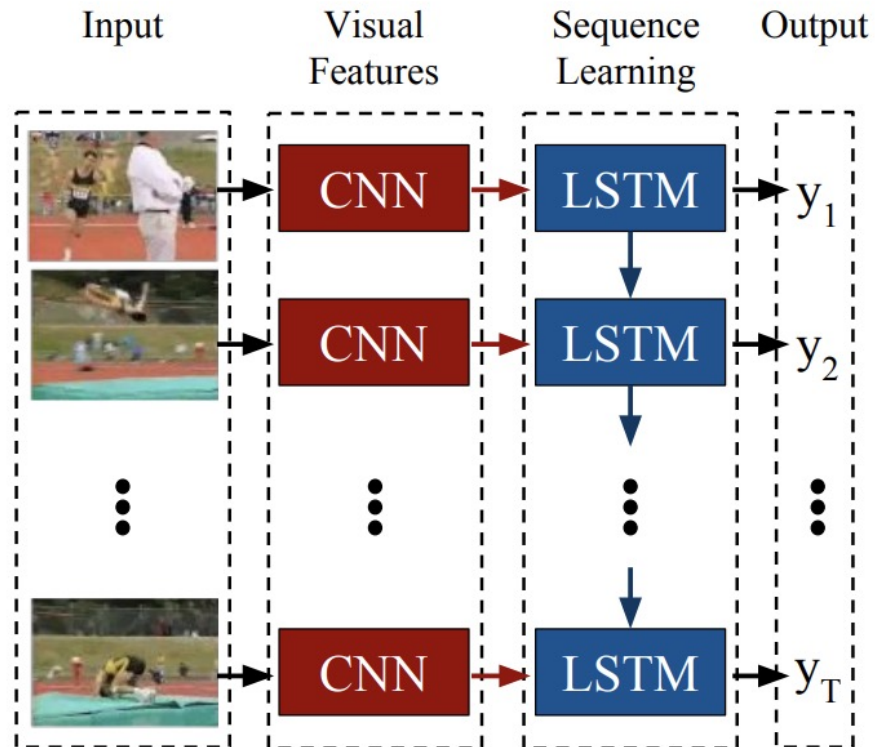
Model: "sequential\_6"

Layer (type)	Output Shape	Param #
conv2d_15 (Conv2D)	(None, 4, 118, 158, 32)	896
conv2d_16 (Conv2D)	(None, 4, 114, 154, 32)	25632
max_pooling3d_10 (MaxPooling3D)	(None, 1, 57, 77, 32)	0
batch_normalization_5 (Batch Normalization)	(None, 1, 57, 77, 32)	128
dropout_10 (Dropout)	(None, 1, 57, 77, 32)	0
conv2d_17 (Conv2D)	(None, 1, 57, 77, 64)	51264
max_pooling3d_11 (MaxPooling3D)	(None, 1, 29, 39, 64)	0
dropout_11 (Dropout)	(None, 1, 29, 39, 64)	0
flatten_5 (Flatten)	(None, 72384)	0
dense_5 (Dense)	(None, 2)	144770

=====  
Total params: 222,690  
Trainable params: 222,626  
Non-trainable params: 64



# Long Term Recurrent CNN



LRCN processes the (possibly) variable-length visual input with a CNN, whose outputs are fed into a stack of recurrent sequence models. A variable-length prediction is then produced.

Both the CNN and LSTM weights are shared across time, resulting in a representation that scales to arbitrarily long sequences.



# Long Term Recurrent CNN



Backhand



Backhand Slice



Backhand2hands

→ Backhand



Forehand Flat



Forehand Open Stance




Forehand Slice

→ Forehand

All serves were combined and labeled as serves. Backhand volley, forehand volley, and smash are kept separate, resulting in 6 classes.

# Long Term Recurrent CNN



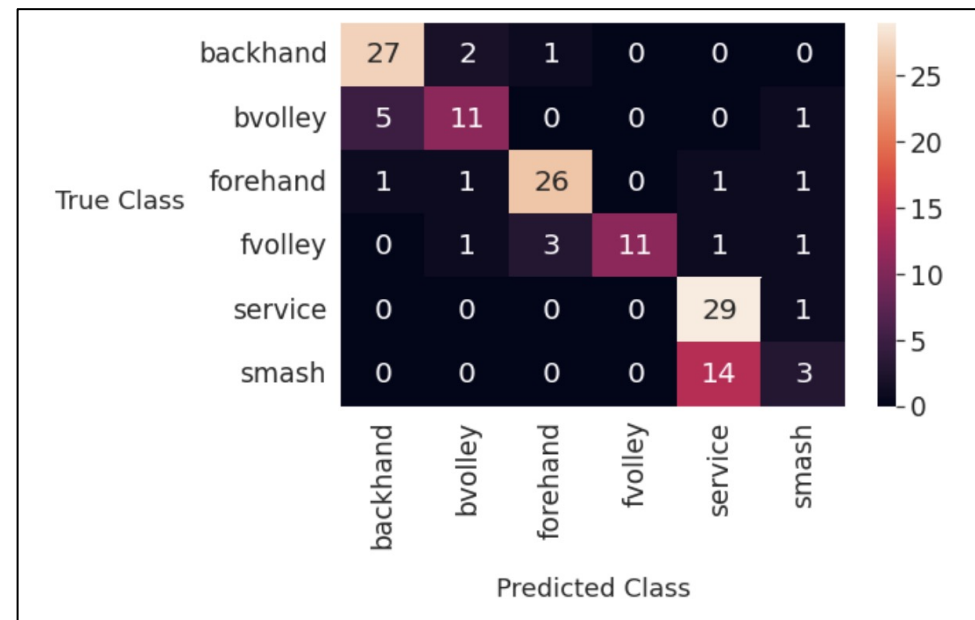
## LSTM Architecture

Layer (type)	Output Shape	Param #
dropout_20 (Dropout)	(None, 16, 2048)	0
lstm_10 (LSTM)	(None, 16, 128)	1114624
time_distributed_20 (TimeDistributed)	(None, 16, 128)	0
time_distributed_21 (TimeDistributed)	(None, 16, 6)	774
lambda_10 (Lambda)	(None, 6)	0
Total params: 1,115,398		
Trainable params: 1,115,398		
Non-trainable params: 0		

# Long Term Recurrent CNN

```
Epoch 300/300
7/8 [=====>....] - ETA: 0s - loss: nan - categorical_accuracy: 0.9592
Epoch 300: saving model to drive/MyDrive/DeepLearning_Test/experiments/base_model/checkpoints/lstm_weights.0300-0.743.hdf5
8/8 [=====] - 0s 52ms/step - loss: nan - categorical_accuracy: 0.9615 - val_loss: 0.7426 - val_categorical_accuracy: 0.7899
```

	precision	recall	f1-score	support
backhand	0.82	0.90	0.86	30
bvolley	0.73	0.65	0.69	17
forehand	0.87	0.87	0.87	30
fvolley	1.00	0.65	0.79	17
service	0.64	0.97	0.77	30
smash	0.43	0.18	0.25	17
accuracy			0.76	141
macro avg	0.75	0.70	0.70	141
weighted avg	0.76	0.76	0.74	141



# Conclusion & Future Work

Model	Number of Output Classes Used	Test Accuracy
Base CNN	12	11%
Binary Classification CNN	2	87%
Binary Classification CNN – 4D input	2	60%
LRCNN	6	76%

- 🎾 Deploy model using Streamlit
- 🎾 Build more complex model in the future and identify all 12 classes
- 🎾 Solve overfitting problem
- 🎾 Solve problem of Nans in training loss of LRCNN
- 🎾 Have model run in real-time over a tennis video for live shot prediction

# References

🟡 Donahue, Jeff, et al. *Long-term Recurrent Convolutional Networks for Visual Recognition and Description*, 2014

🟡 <https://github.com/ganasank/CS230>

🟡 [https://github.com/chow-vincent/tennis\\_action\\_recognition](https://github.com/chow-vincent/tennis_action_recognition)

🟡 <http://thetis.image.ece.ntua.gr/textfiles/thetis.pdf>

🟡 <http://thetis.image.ece.ntua.gr/>